



Multiprocessors

Erik Hagersten
Uppsala University

Outline of these lectures

1. Uniprocessors
2. Multiprocessors
 - ✱ Coherence
 - ✱ Memory ordering
 - ✱ Vector instructions, GPUs
3. Multicores & Manycores
4. Optimizing for speed

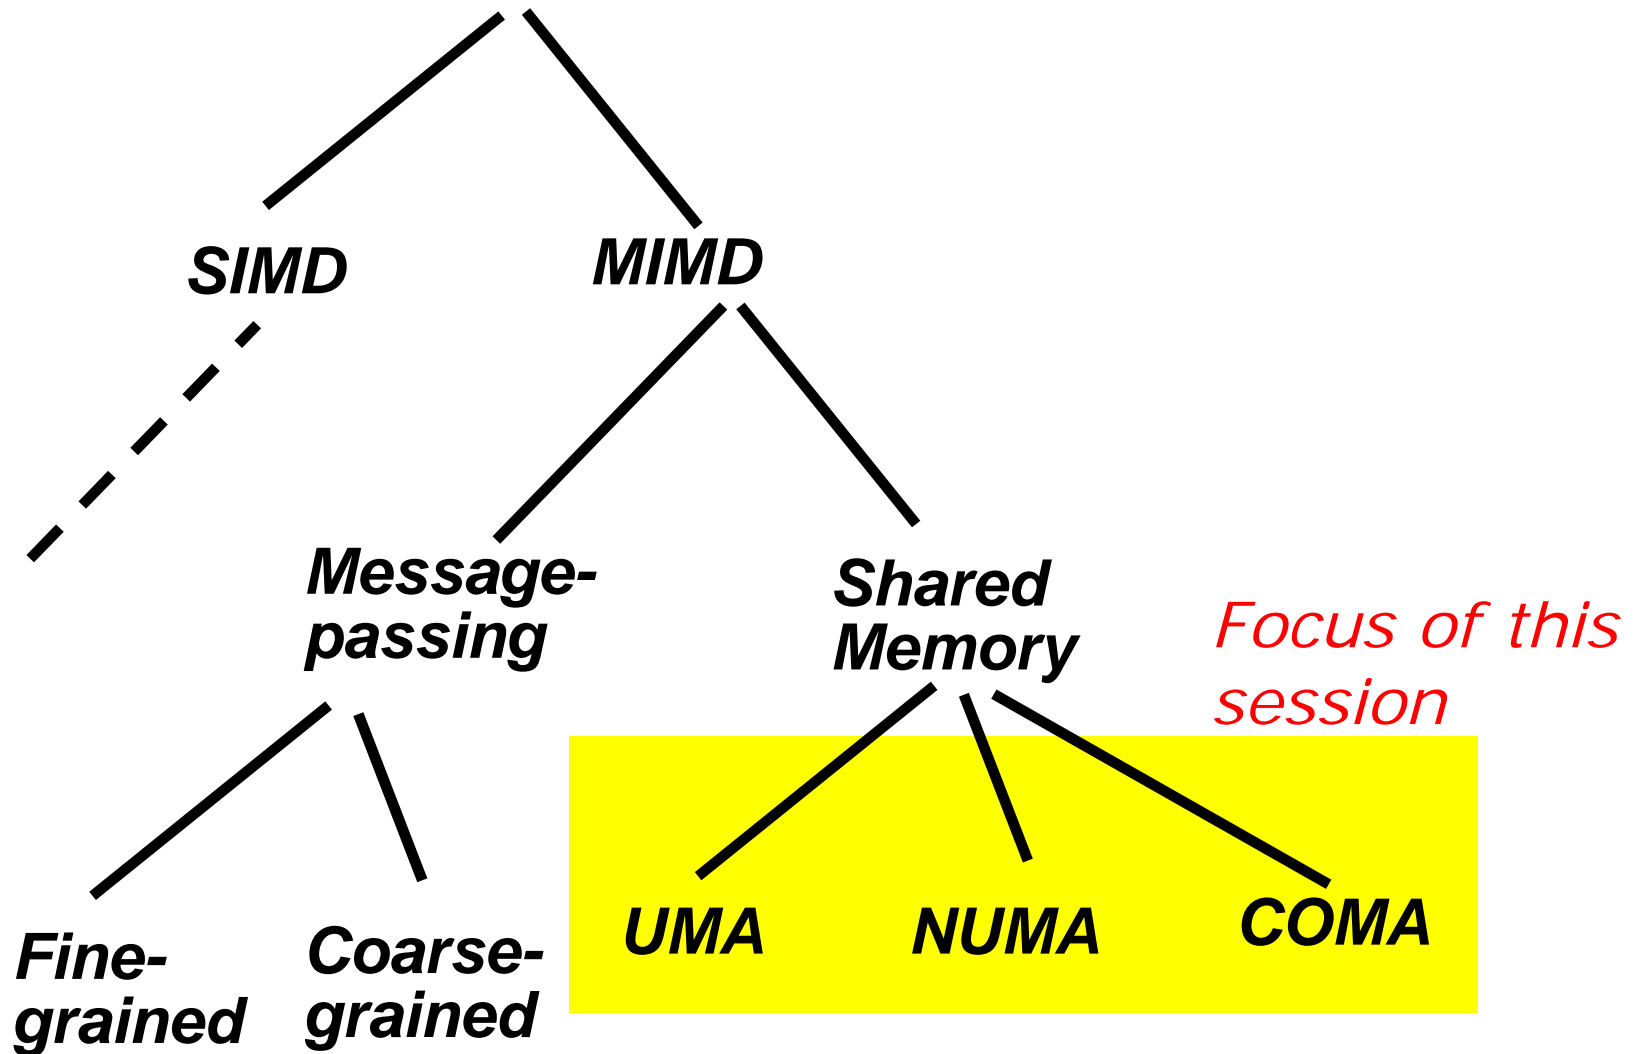
The era of the “supercomputer” multiprocessors

- The one with the most blinking lights wins
- The one with the strangest languages wins
- The niftier the better!





Taxomy for Architectures [Flynn]

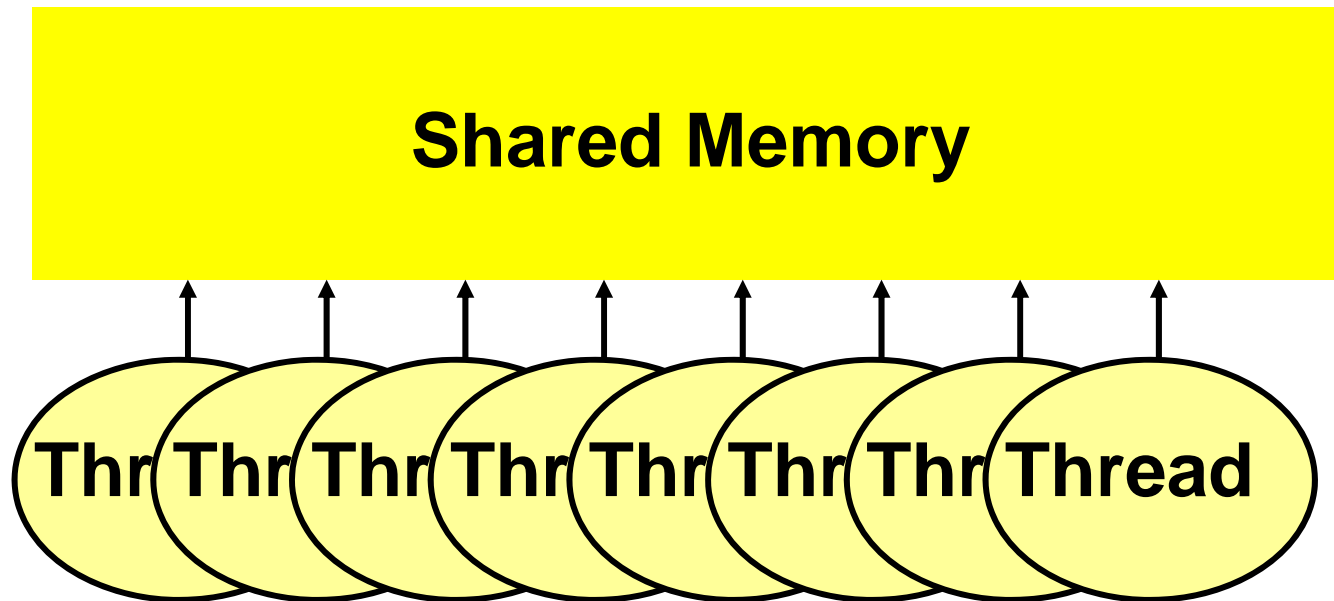




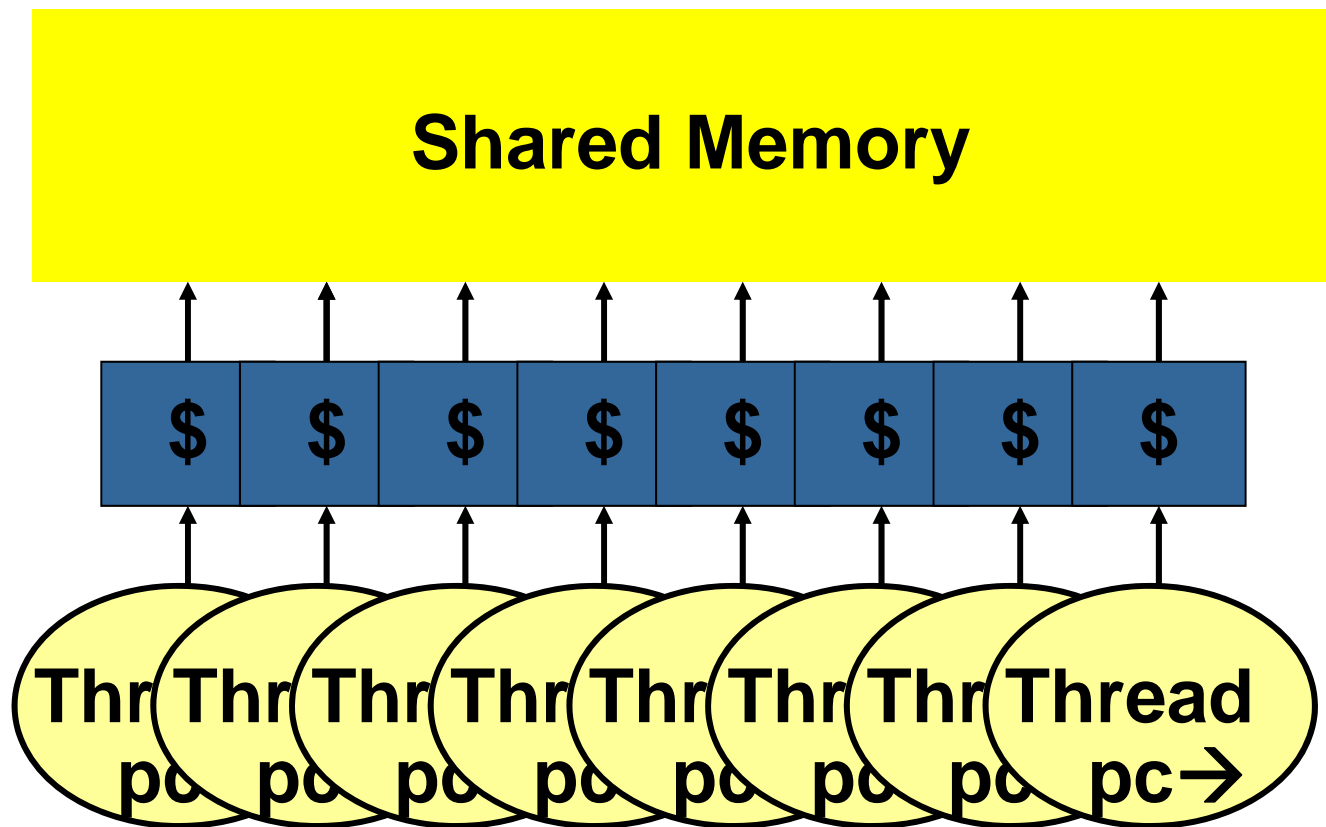
Coherent Shared Memory

Erik Hagersten
Uppsala University

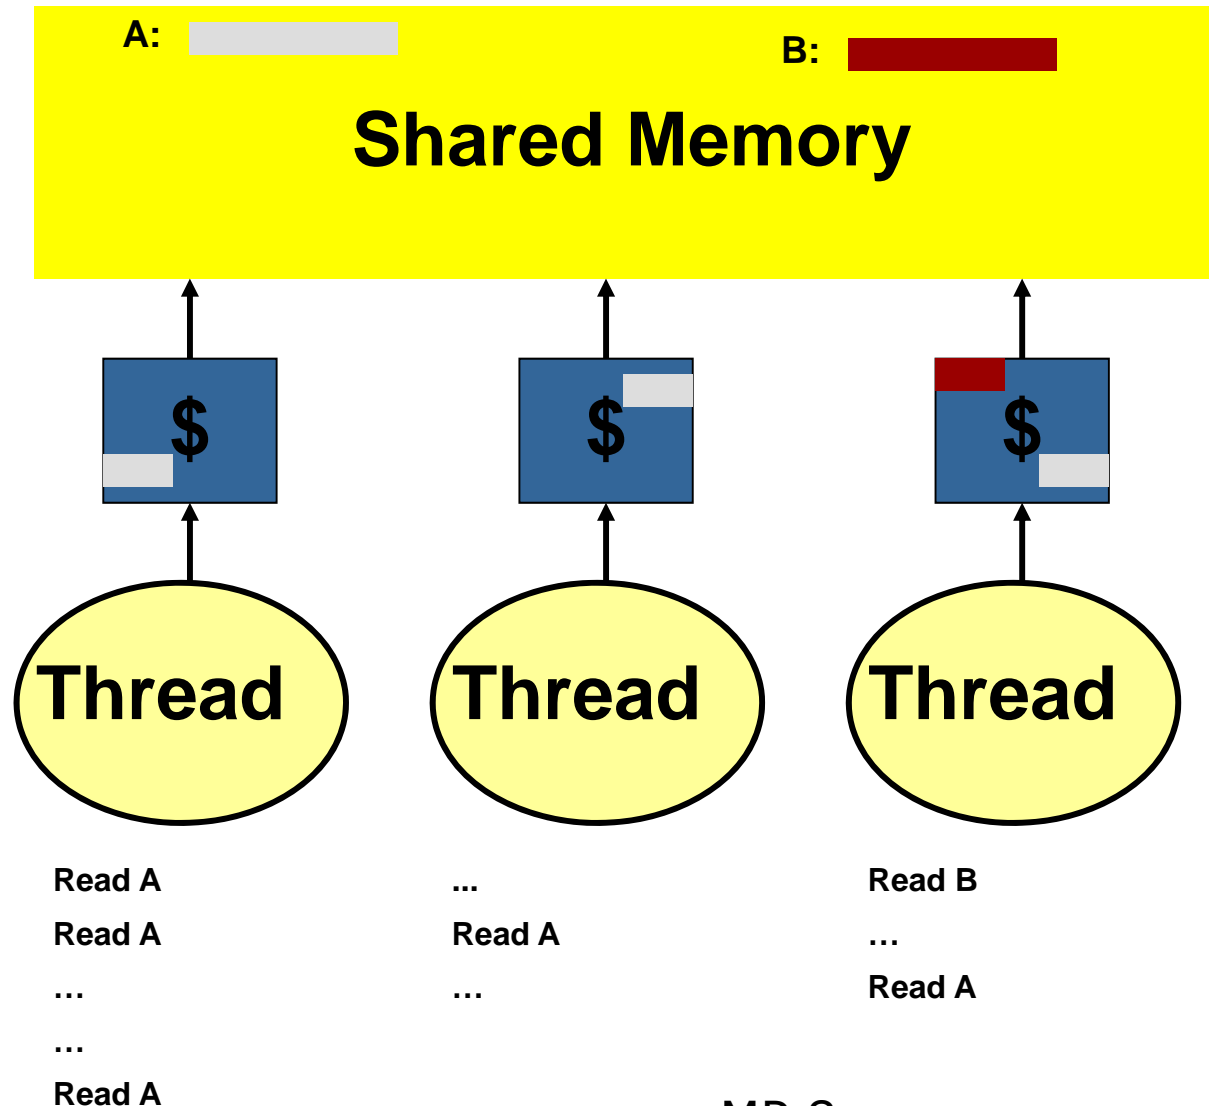
Programming Model:



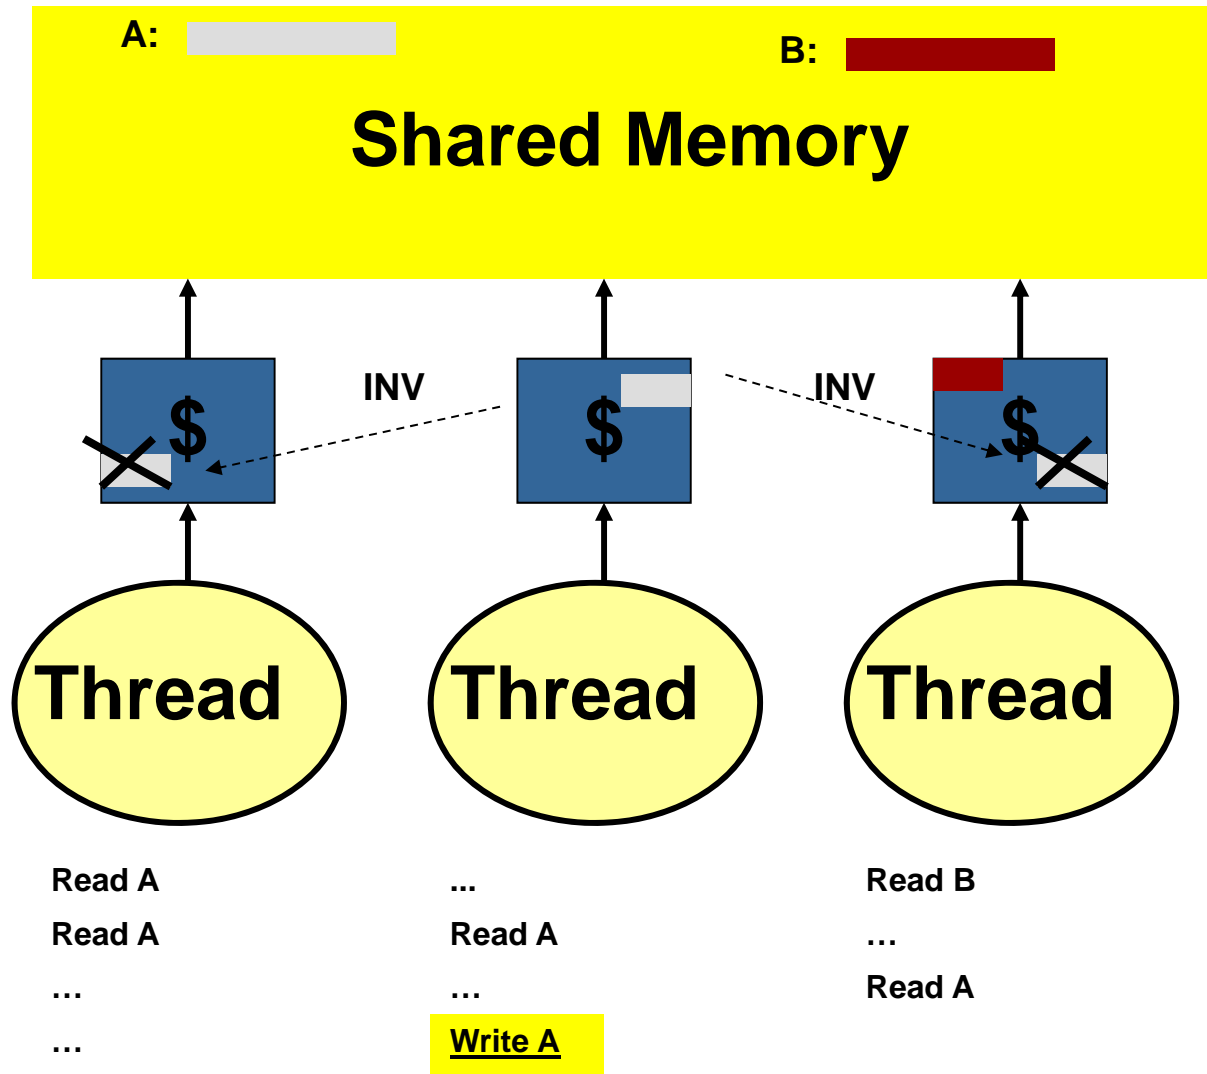
Adding Caches: More Concurrency



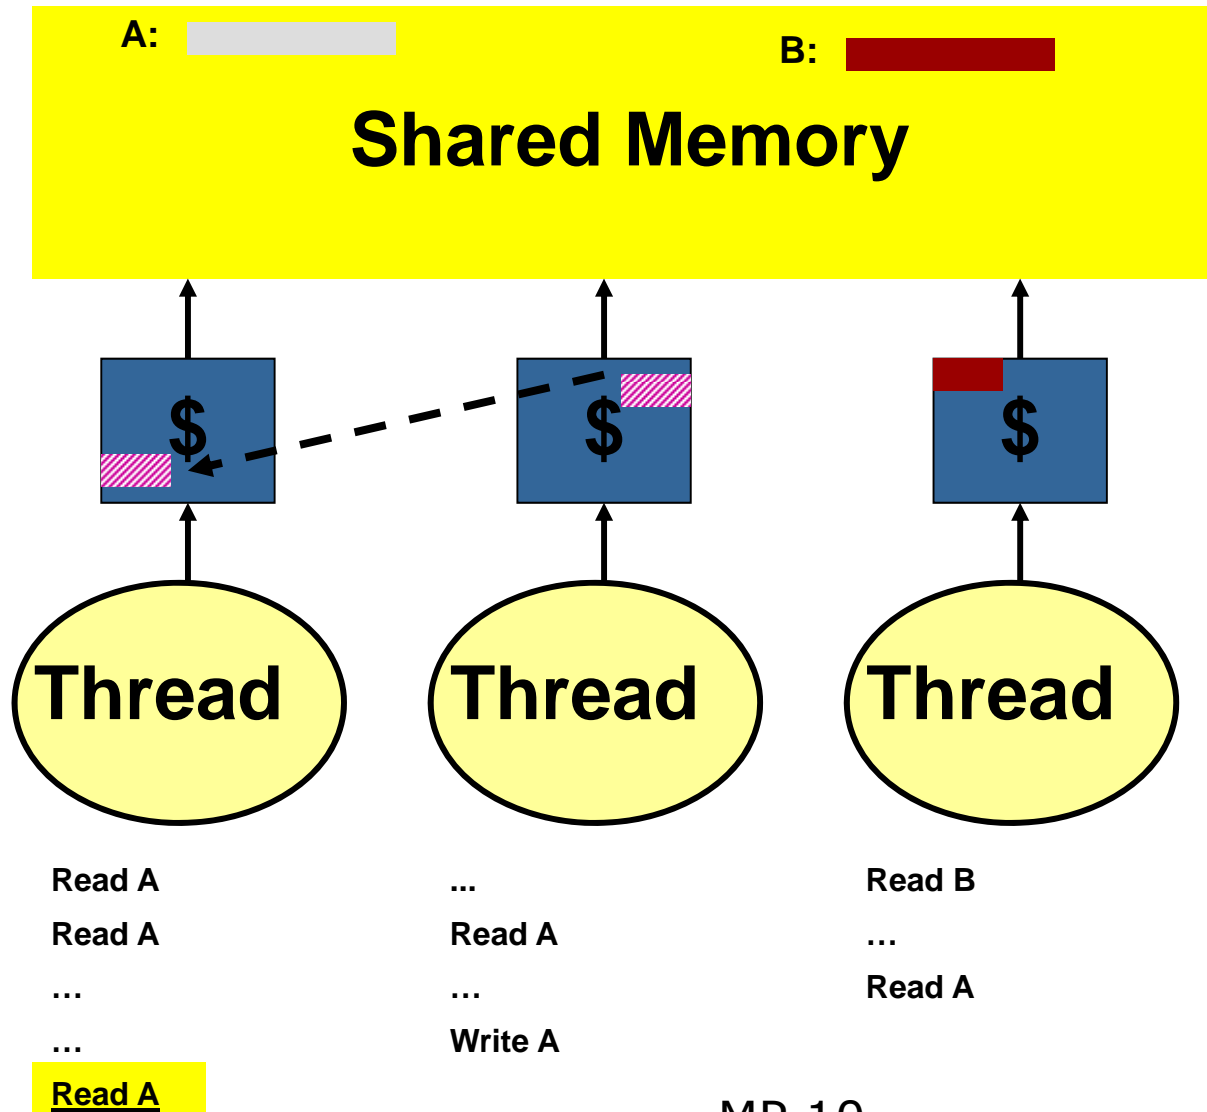
Caches: Automatic Replication of Data



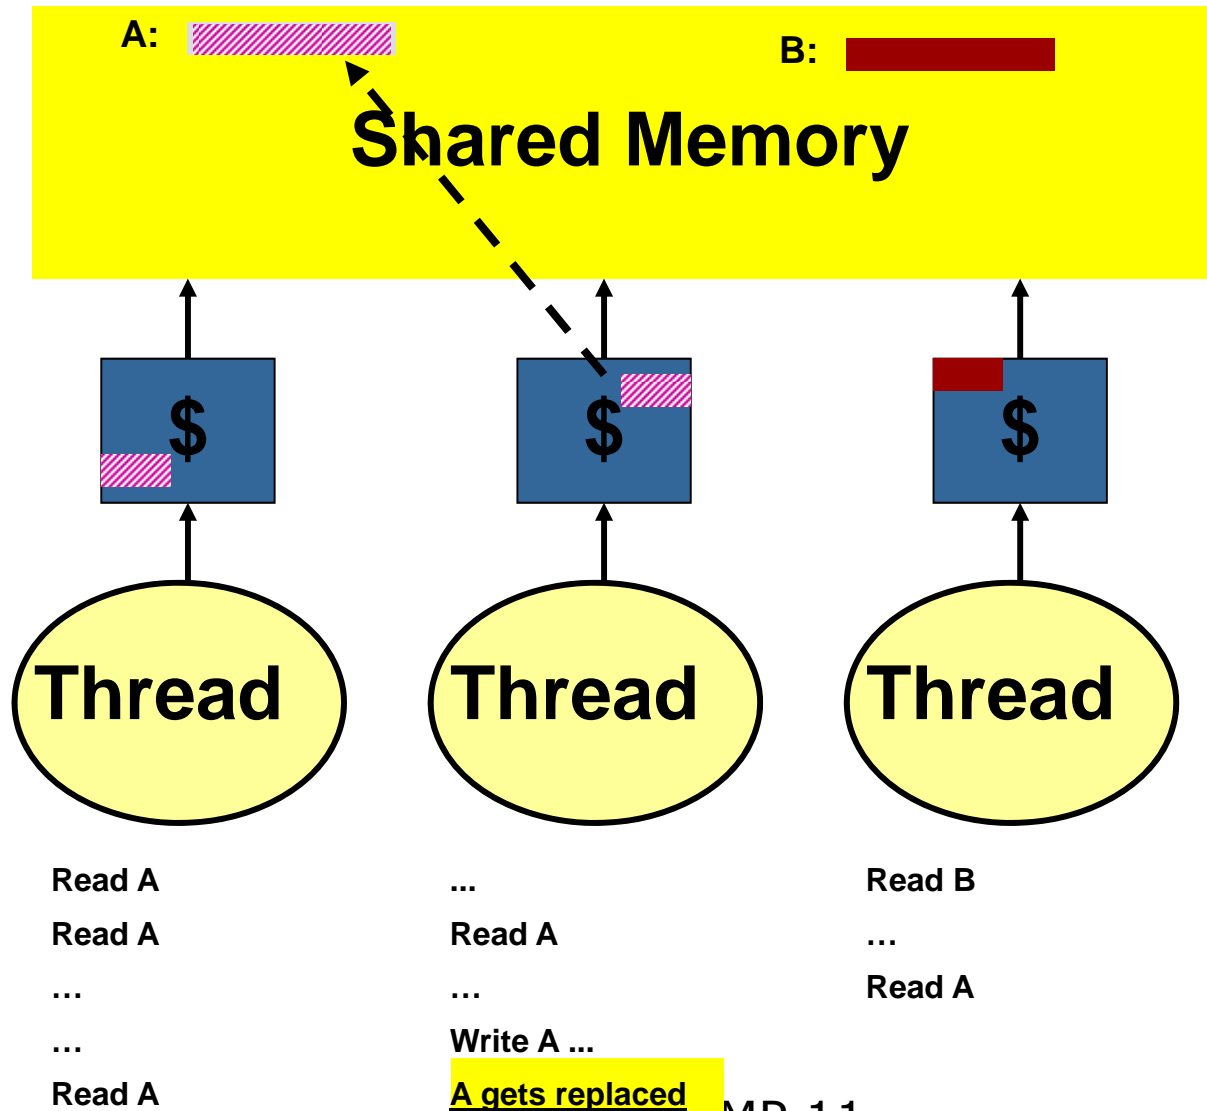
The Cache Coherent Memory System



The Cache Coherent \$2\$



Writeback



Summing up Coherence

Sloppy: there can be many copies of a datum, but only one value

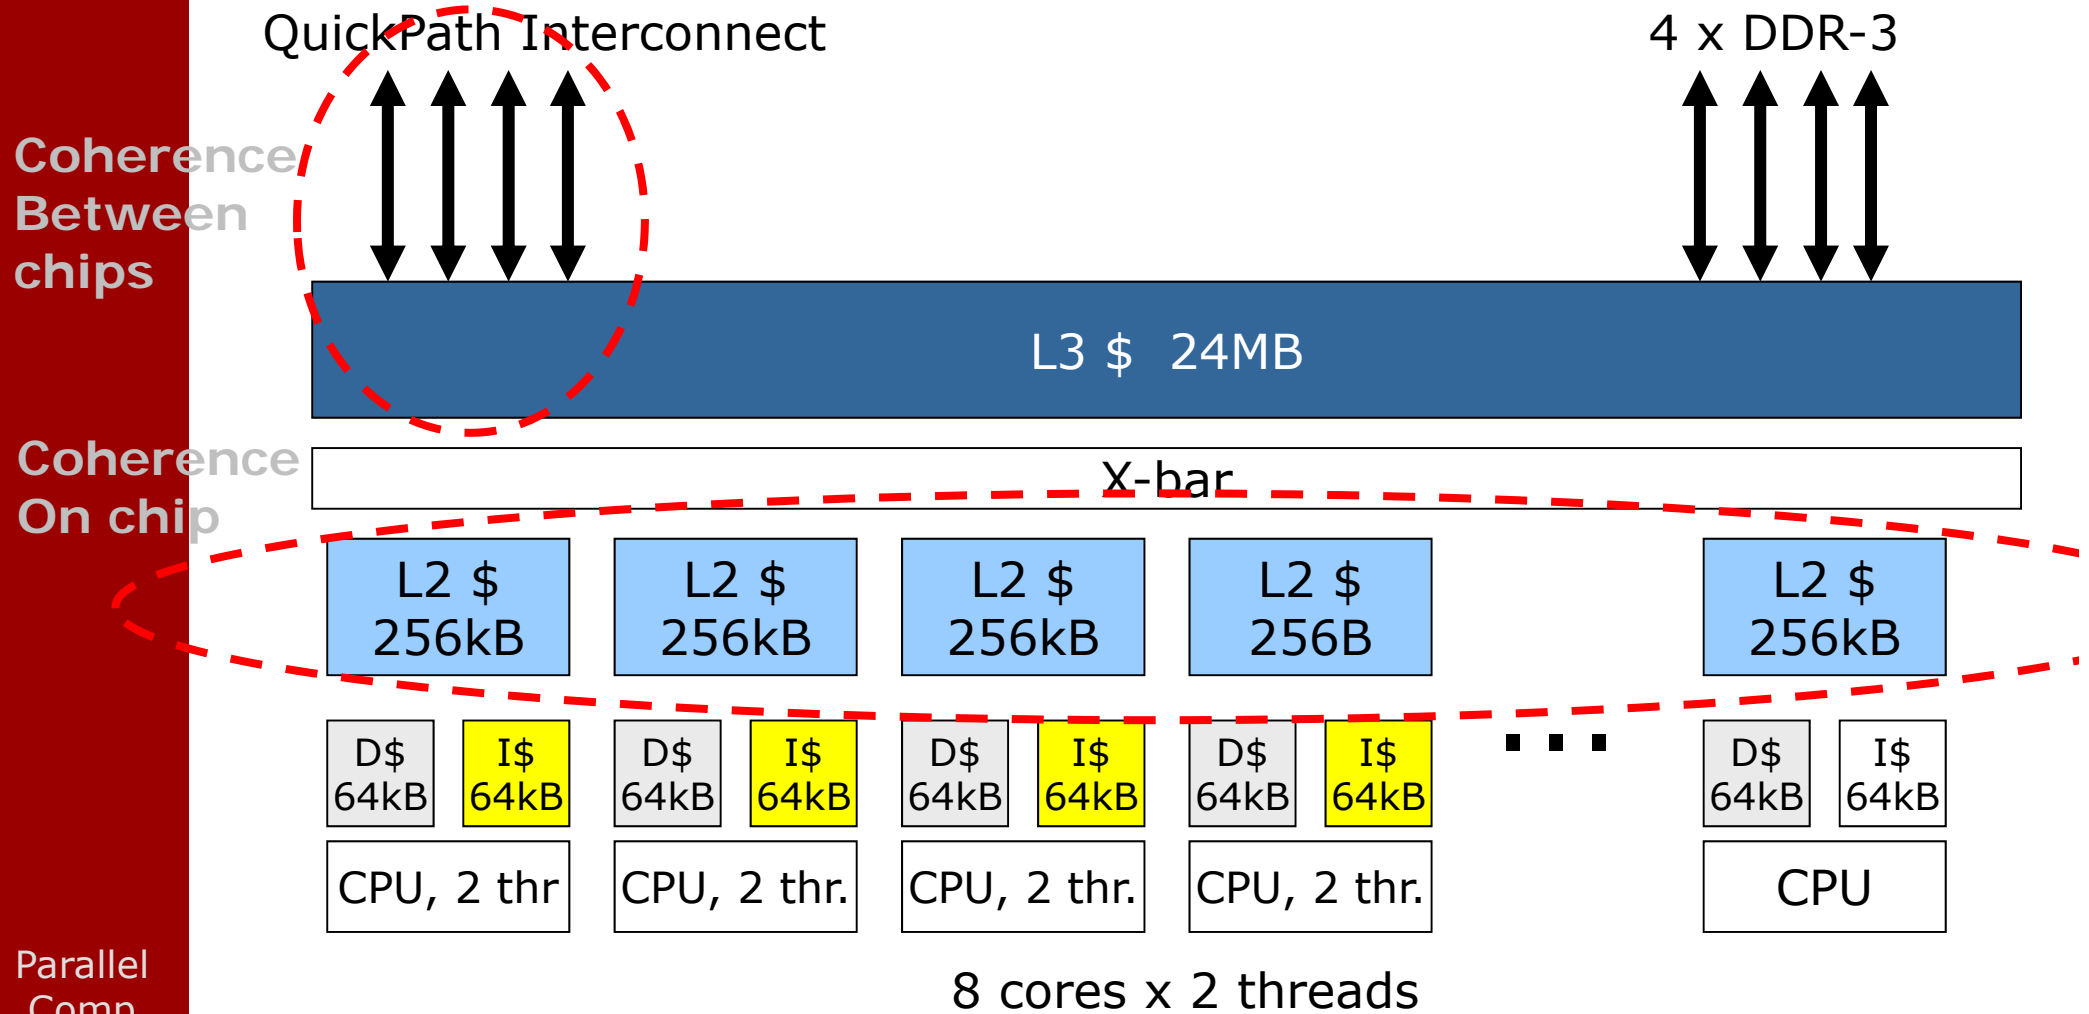
Too strong definition!

Coherence: *There is a single global order of value changes to each datum*

Memory order/model: *Defines the order between accesses to many data*



Where does coherence matter?



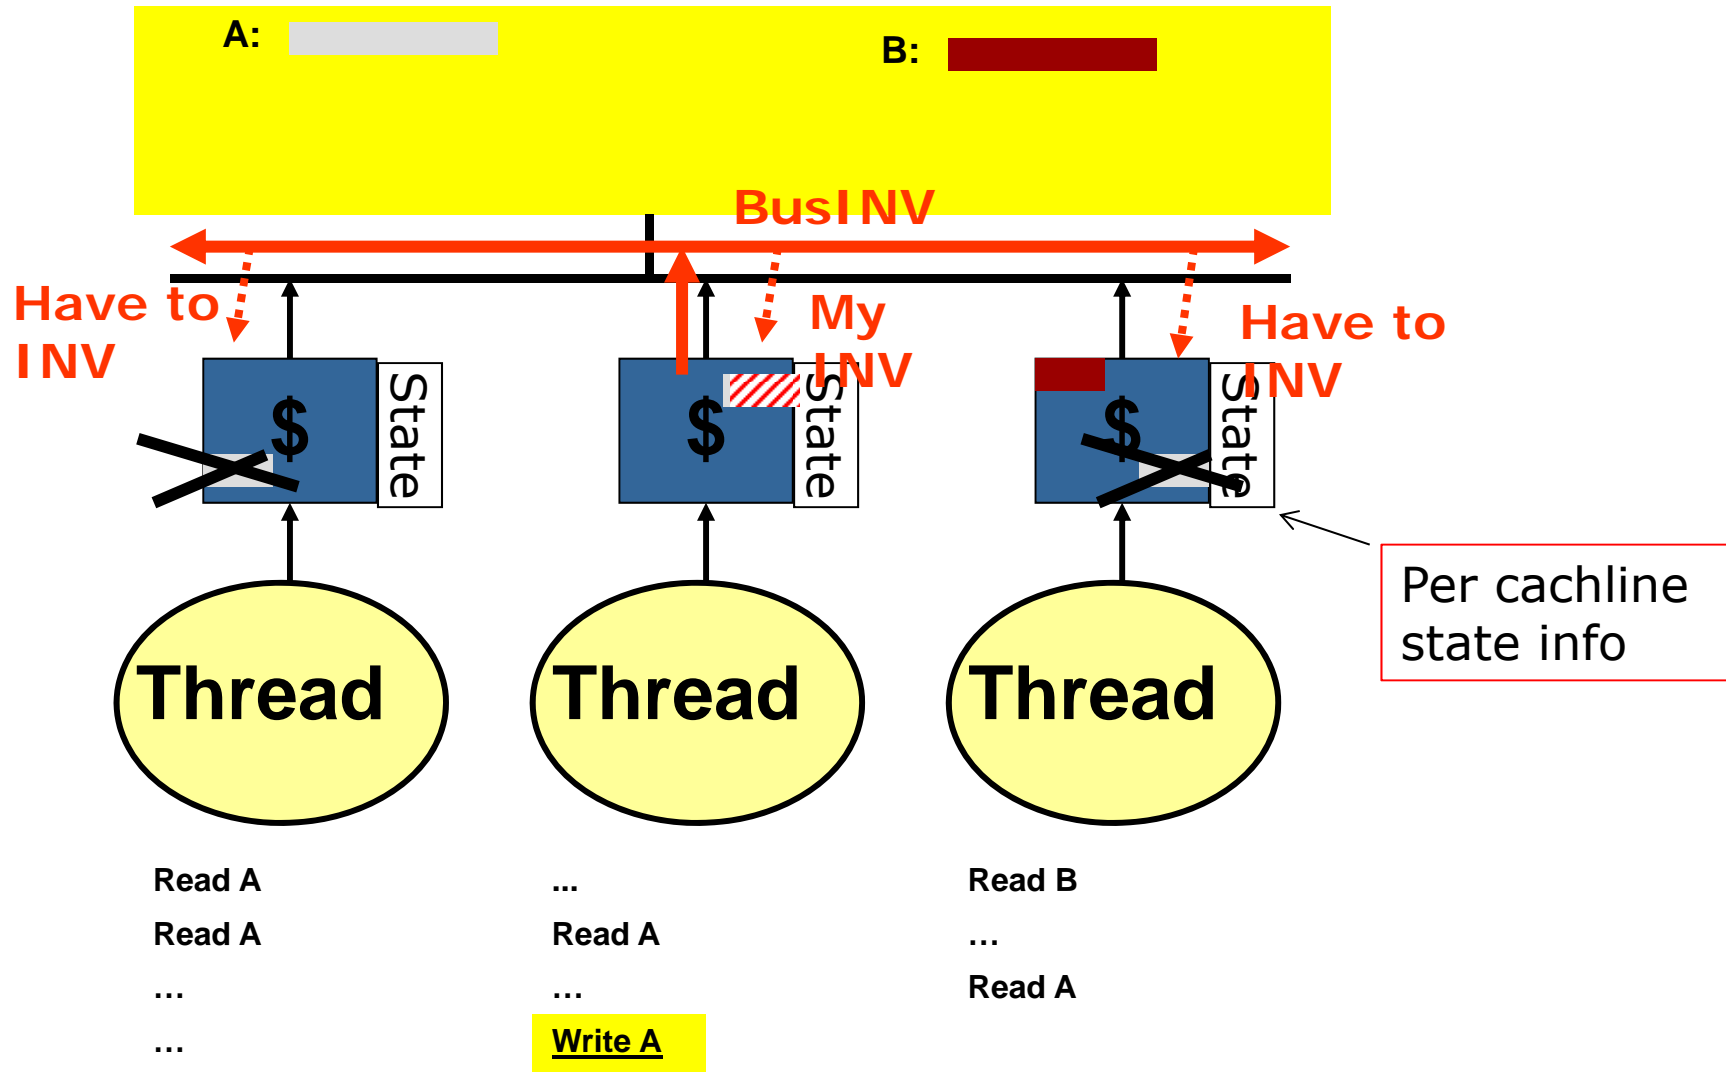


UPPSALA
UNIVERSITET

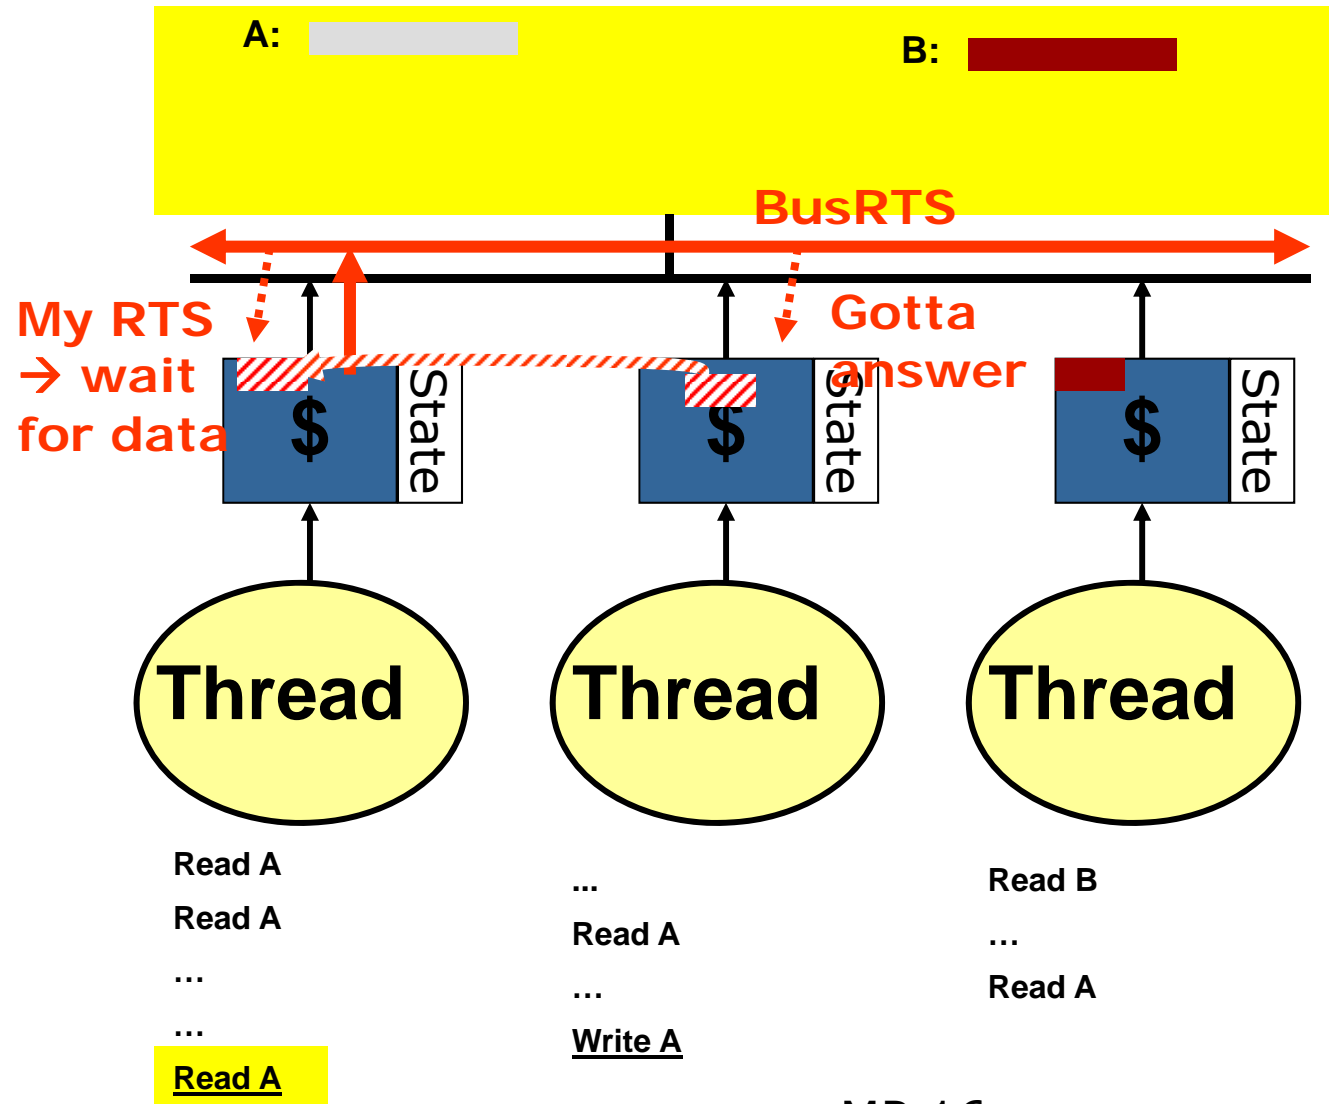
Implementing Coherence

Parallel
Comp
2012

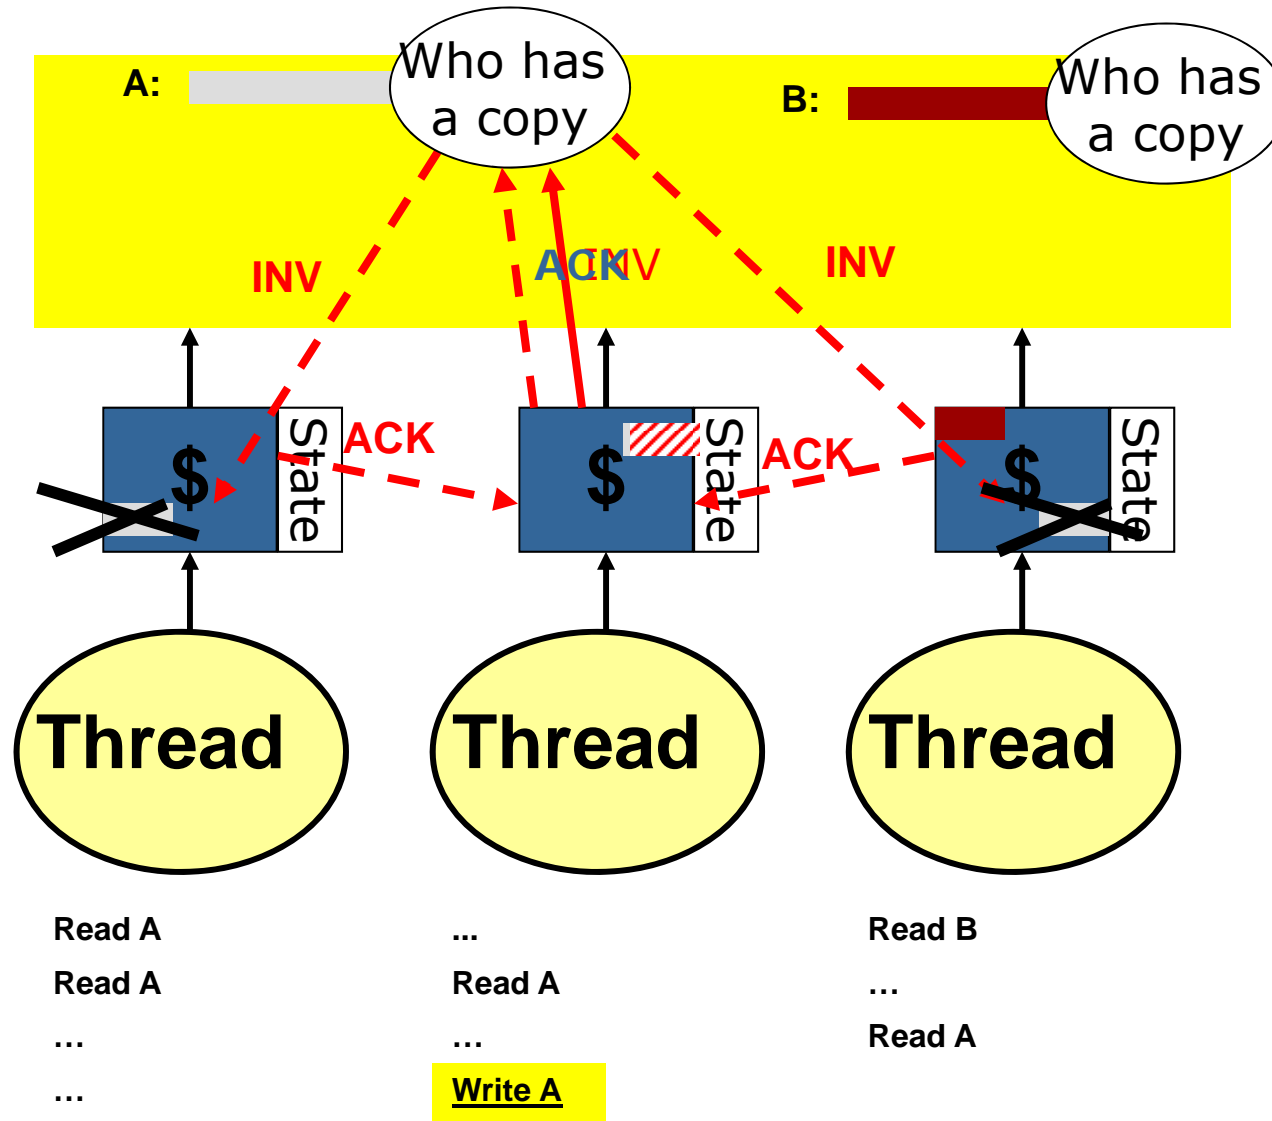
"Upgrade" in snoop-based



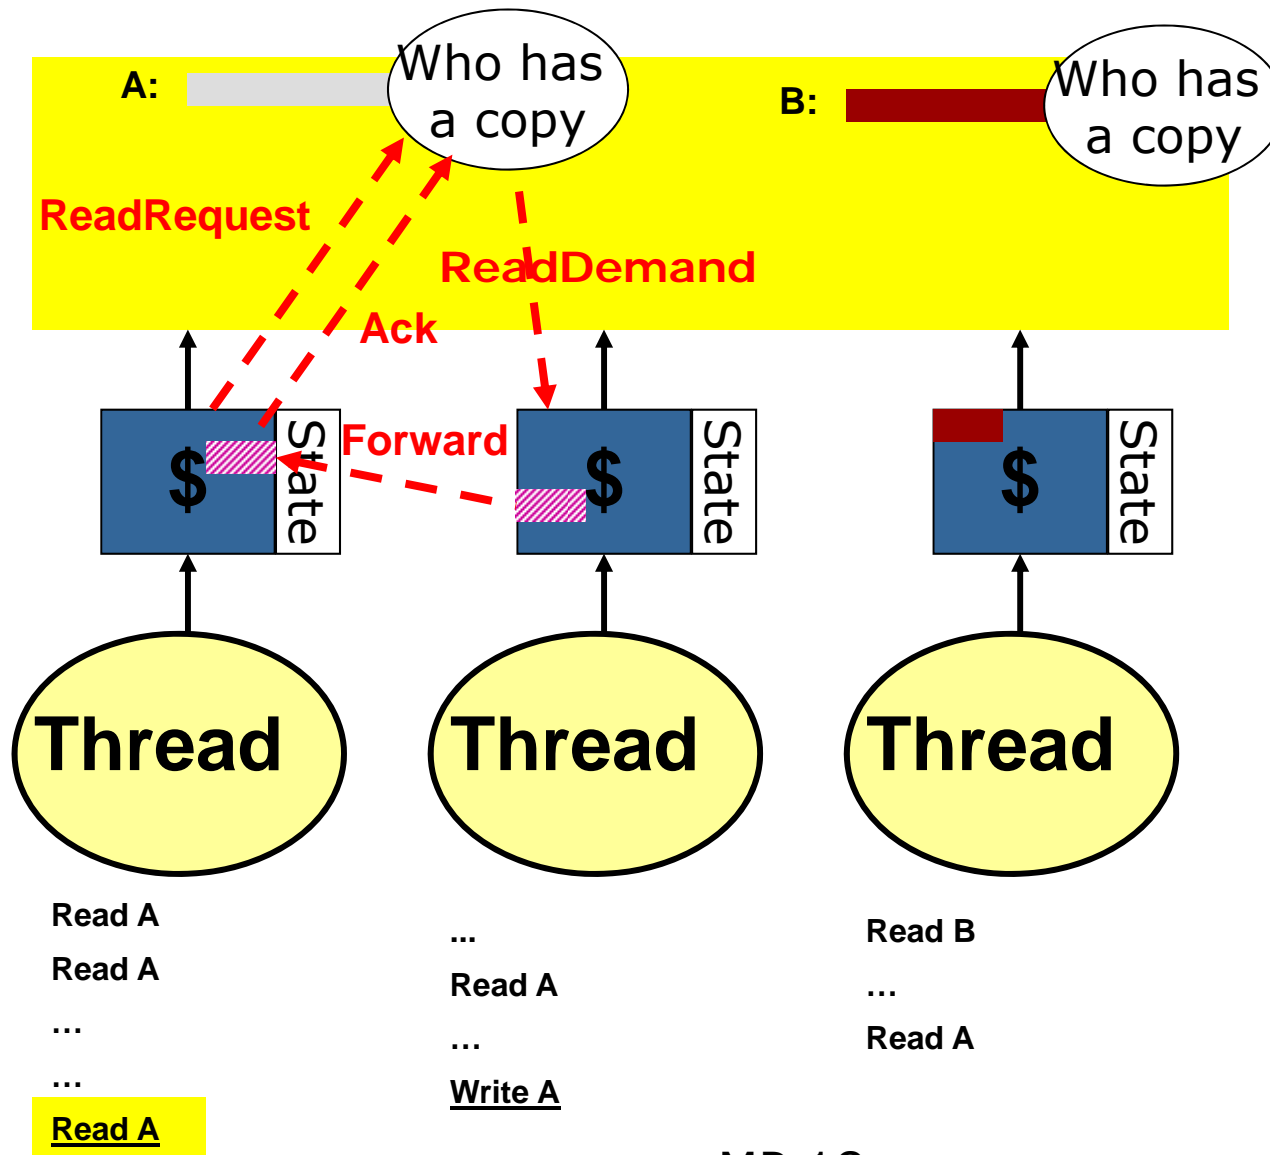
Cache-to-cache in snoop-based



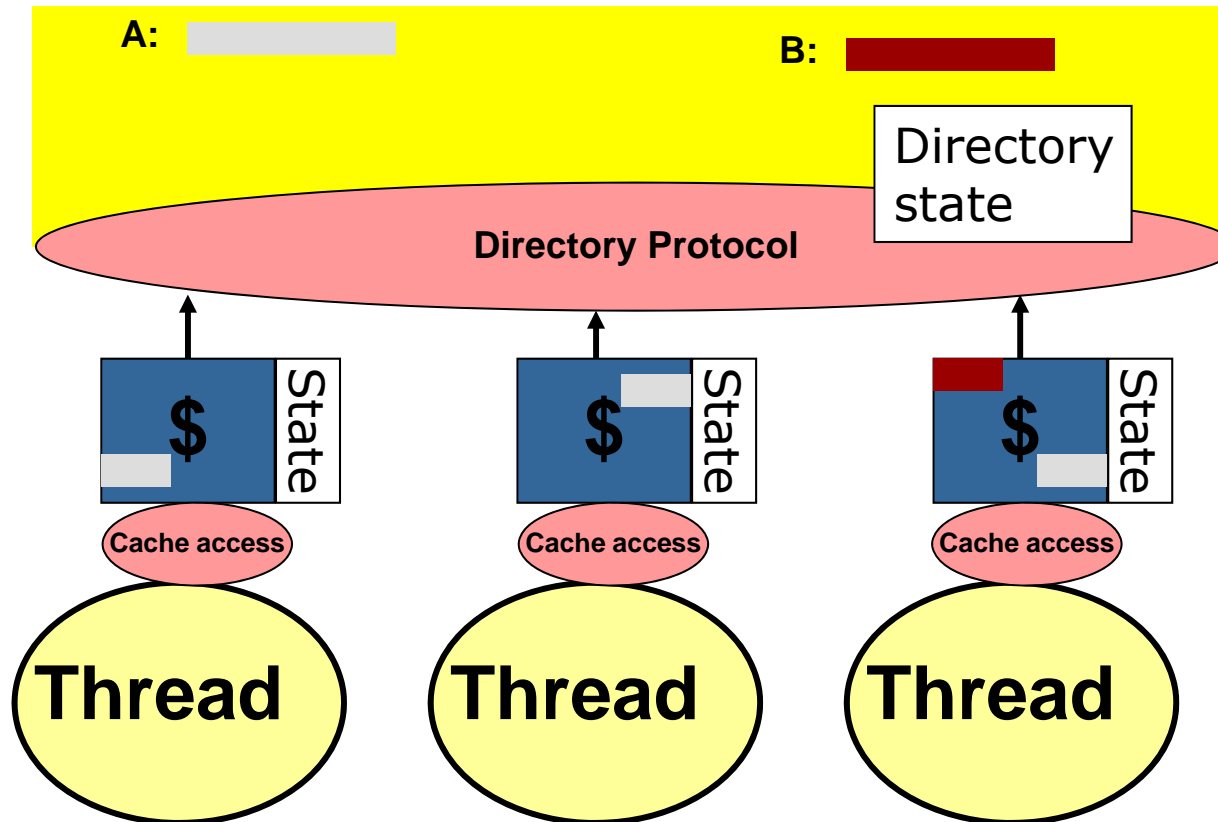
"Upgrade" in dir-based




Cache-to-cache in dir-based



Directory-based coherence: Per-cacheline info in the memory

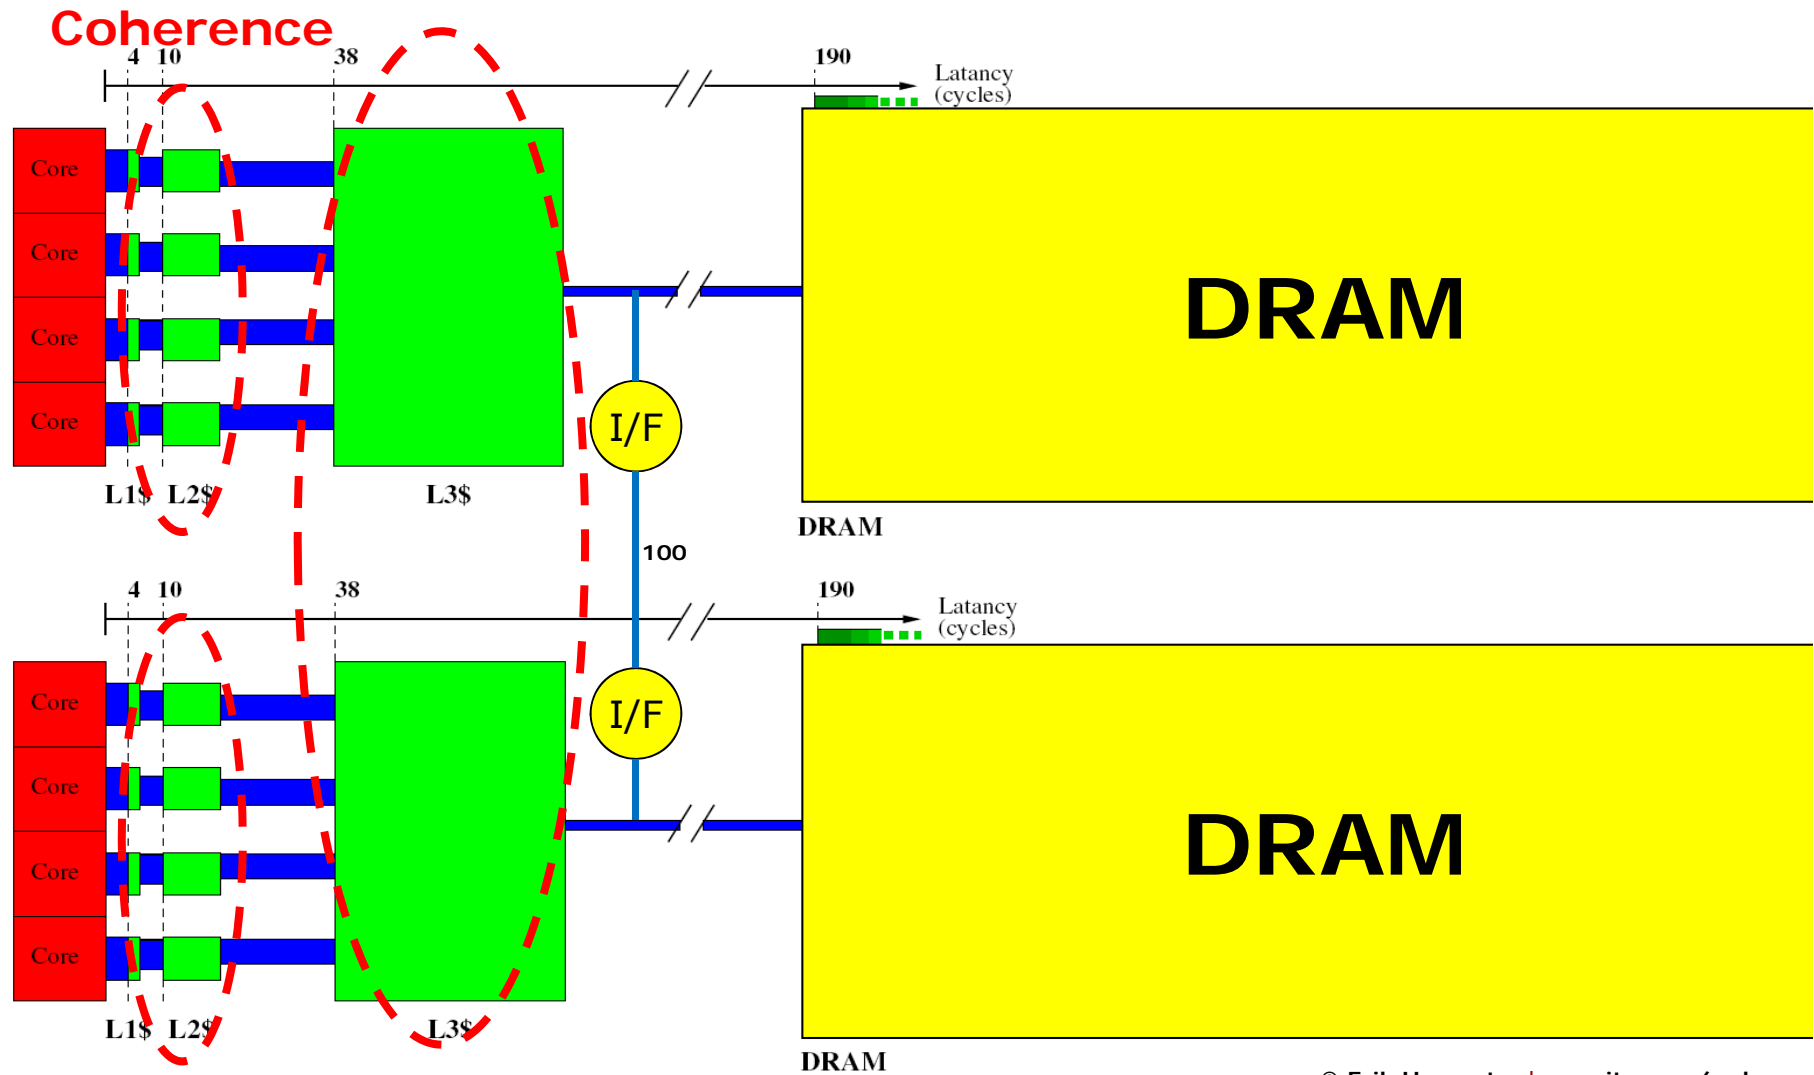


Directory-based snooping: NUMA. Per-cacheline info in the home node

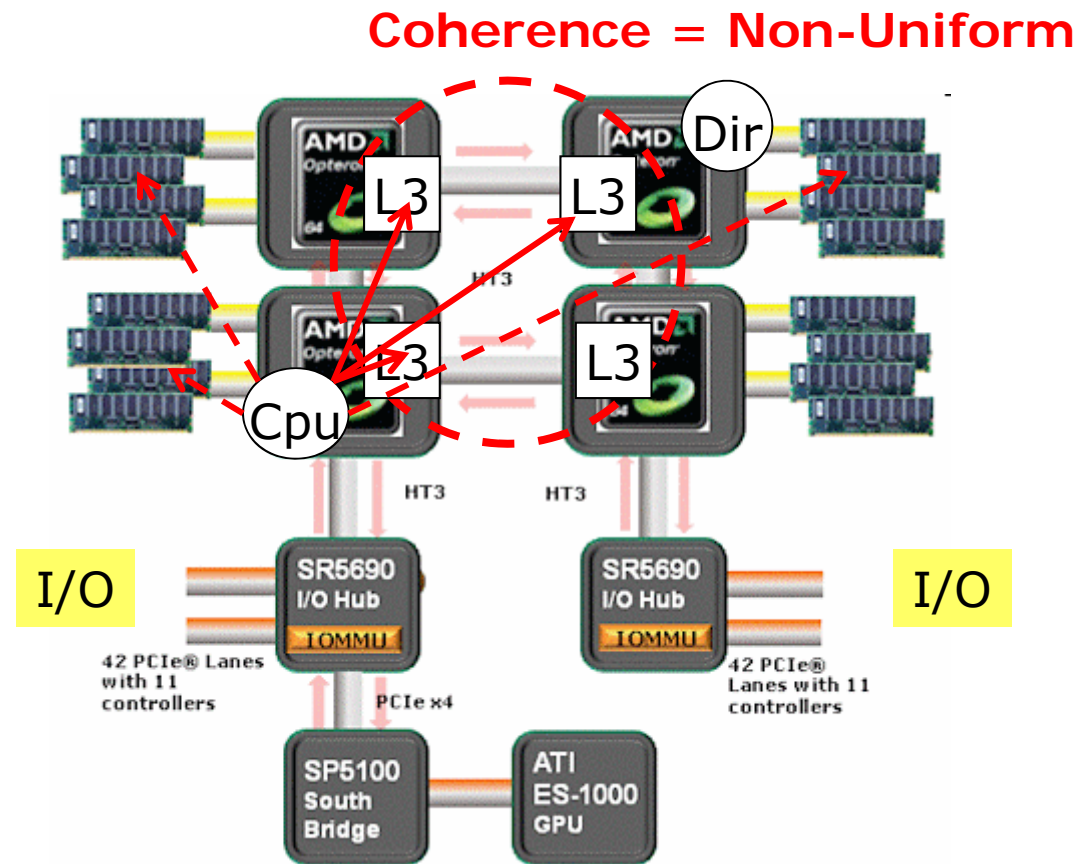


Multisocket

Coherence = Non-Uniform



AMD Multi-socket Architecture (same applies to Intel multi-sockets)

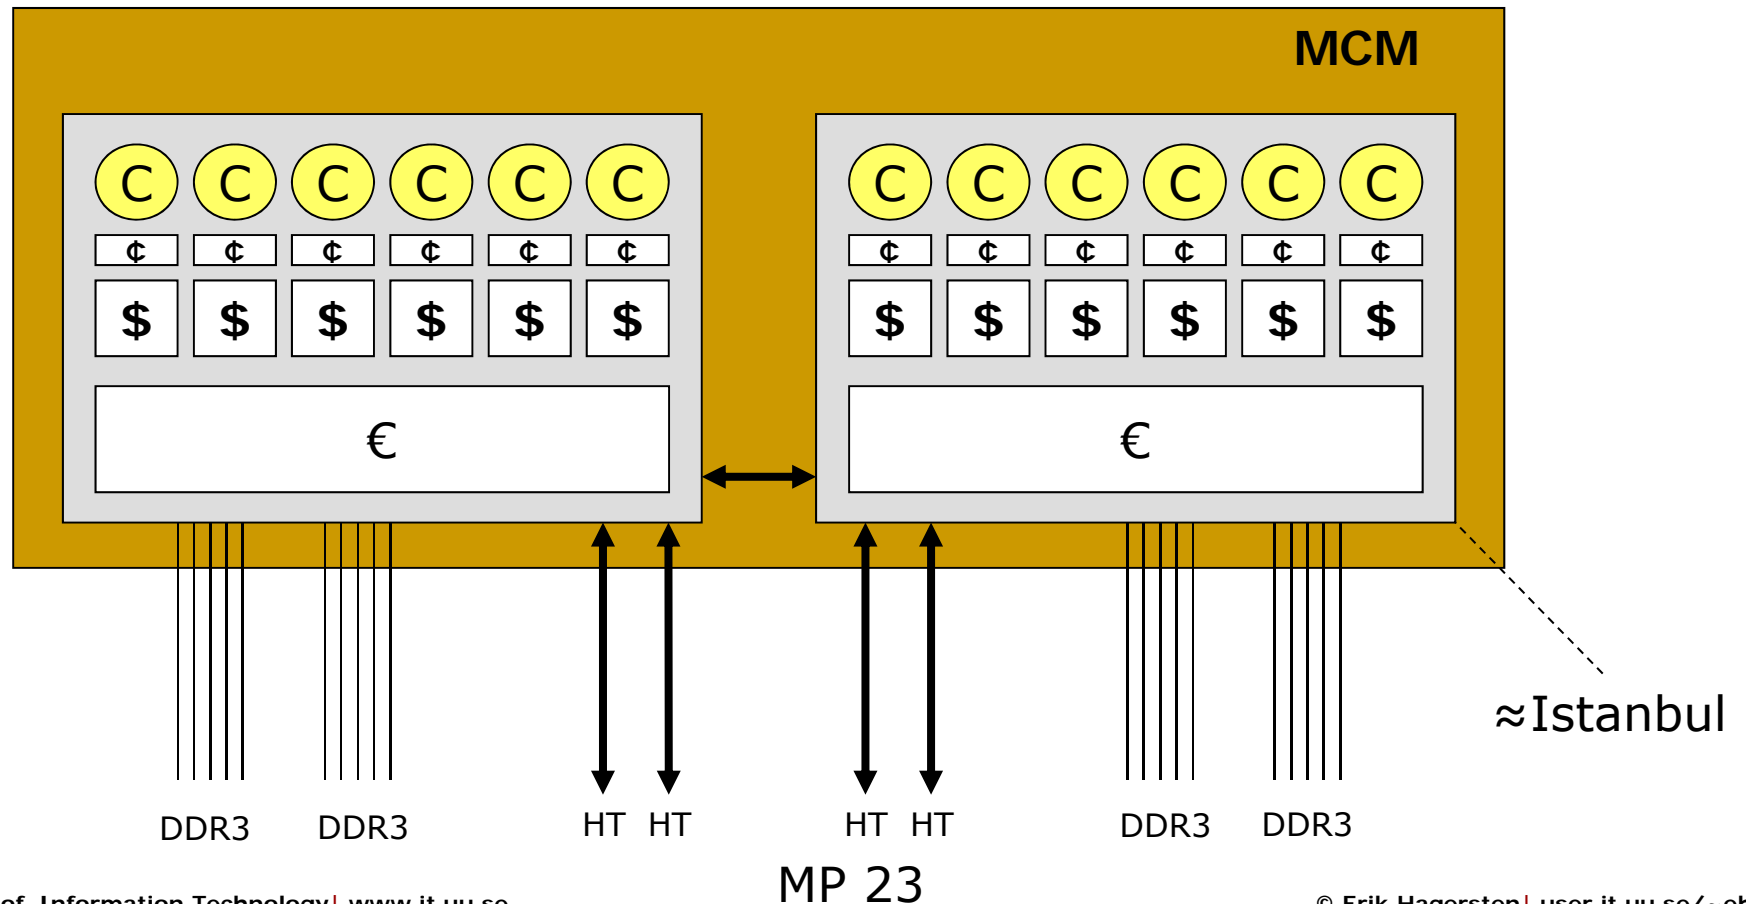


AMD Magny-Cours


NUMA & NUCA on a socket

Non-Uniform Memory Architecture


Non-Uniform Communication Architecture



False sharing: Coherence is maintained with a cache-line granularity



Communication misses even though
the threads do not share data
"the cache line is too large"



Read A
Write A
...
...
Read A

Read E
...
Write E

More Cache Lingo

- **Capacity miss** – too small cache
- **Conflict miss** – limited associativity
- **Compulsory miss** – accessing data the first time
- **Coherence miss** – I would have had the data unless it had been invalidated by someone else
- **Upgrade miss** (only for writes) – I would have had a writable copy, but gave away readable data and downgraded myself to read-only
- **False sharing**: Coherence/downgrade is caused by a shared cacheline, to by shared data:

**False sharing
example:**

Read A
...
Write A
...
Read A

...
Read D
...
Write D

cacheline:


A, B, C, D



Memory Ordering (aka Memory Consistency) -- tricky but important stuff


Erik Hagersten
Uppsala University
Sweden

The Shared Memory Programming Model (Pthreads/OpenMP, ...)




Memory Ordering

- Coherence defines a per-datum valuechange order
- Memory model defines the valuechange order for all the data.



Q: What value will get printed?

Dekker's Algorithm



Q: Is it possible that both A and B win?


Memory Ordering

- Defines the guaranteed memory ordering: *If a thread has seen that A happens before B, what order can the other threads observe?*
- Is a "contract" between the HW and SW guys
- Without it, you can not say much about the result of a parallel execution


Human intuition: There is one global order!

(A' denotes a modified value to the data at addr A)


Thread 1



Thread 2




(LD A happen before ST A')



Human intuition: There is one global order!


(A' denotes a modified value to the data at addr A)

Thread 1



(LD A happend before ST A')


Thread 2




One possible observed order

Another possible observed order

Thread 1 Thread 2




Thread 1 Thread 2



"The intuitive memory order"


Sequential Consistency (Lamport)



- ✱ Global order achieved by *interleaving* all memory accesses from different threads
- ✱ "Programmer's intuition is maintained"
 - Store causality? Yes
 - Does Dekker work? Yes
- ✱ Unnecessarily restrictive ==> performance penalty

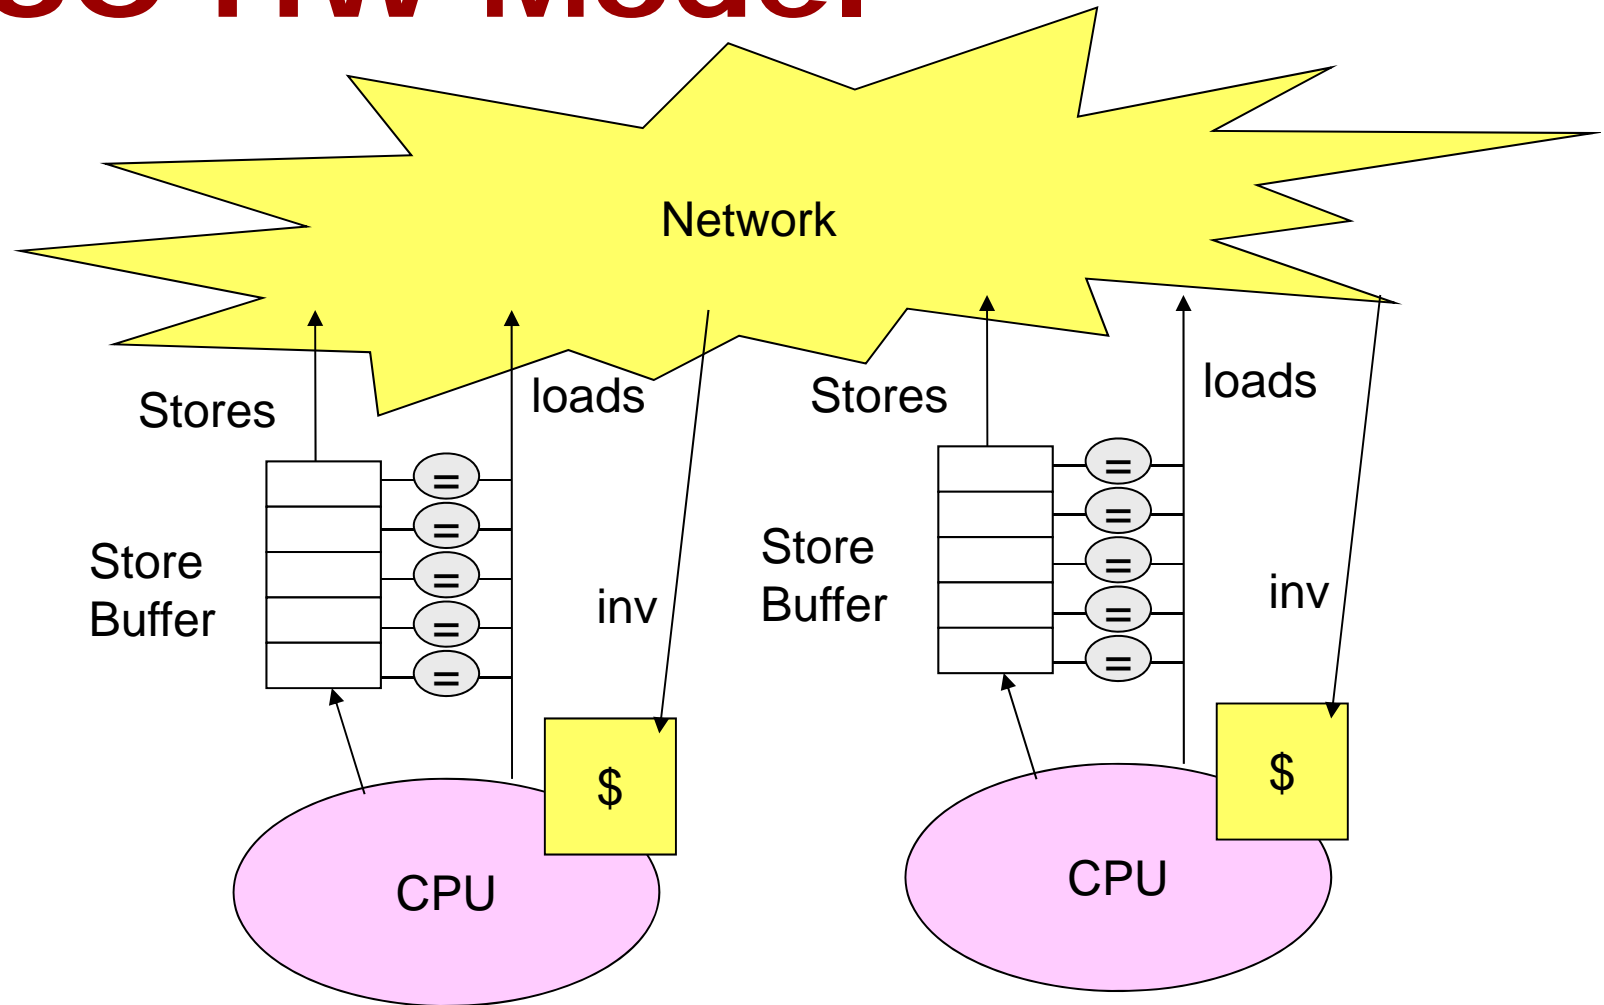
"Almost intuitive memory model"

Total Store Ordering [TSO] (P. Sindhu)



- ✱ Global *interleaving* [order] for all stores from different threads (own stores excepted)
- ✱ Is programmer's intuition is maintained?
 - Store causality? Yes
 - Does Dekker work? No
- ✱ Unnecessarily restrictive ==> performance penalty

TSO HW Model



→ Stores are moved off the critical path


TSO

■ Flag synchronization works

A := data
flag := 1

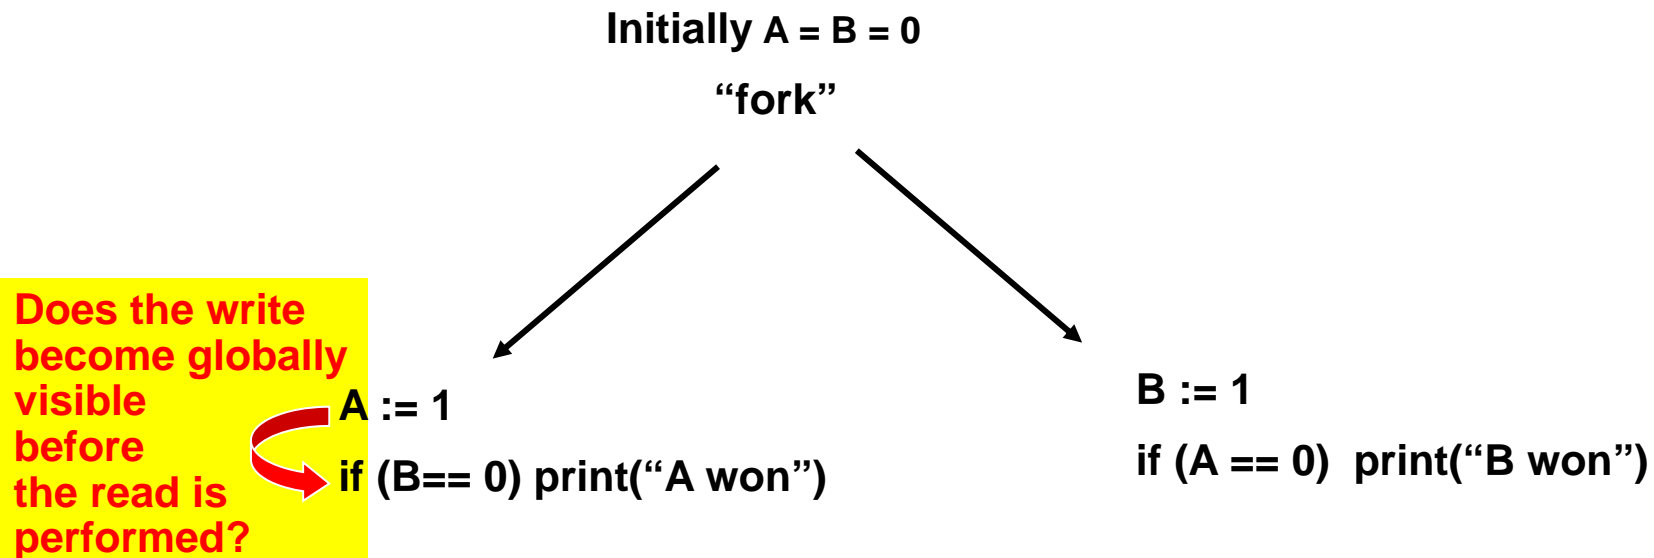
while (flag != 1) {};
X := A

■ Provides causal correctness



Q: What
value will
get printed?
Answer: 1

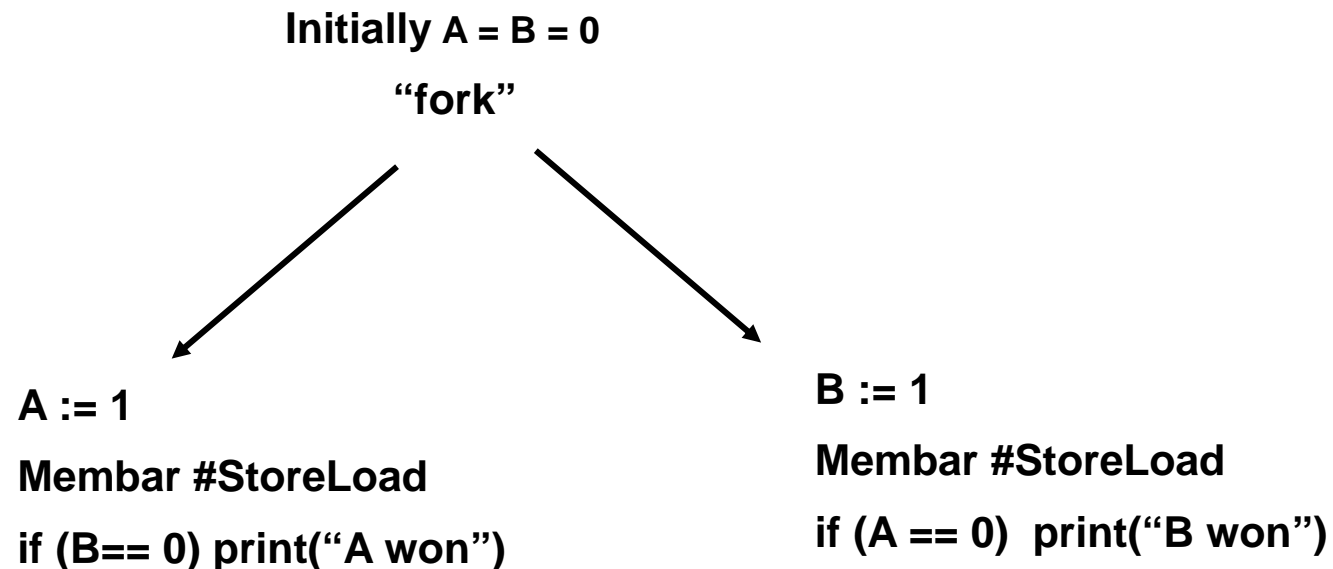
Dekker's Algorithm, TSO



Q: Is it possible that both A and B wins?

Left: The read (i.e., test if $B == 0$) can bypass the store ($A := 1$)
Right: The read (i.e., test if $A == 0$) can bypass the store ($B := 1$)
→ both loads can be performed before any of the stores
→ yes, it is possible that both wins
→ → Dekker's algorithm breaks

Dekker's Algorithm for TSO




Q: Is it possible that both A and B win?

Membar: The read is started after all previous stores have been "globally ordered"

→ behaves like SC

→ Dekker's algorithm works!

Weak/release Consistency (M. Dubois, K. Gharachorloo)



- Most accesses are unordered
 - Programmer's intuition is not maintained?
 - Store causality? No
 - Does Dekker work? No
 - Global order only established when the programmer explicitly inserts memory barrier instructions
- ++ Better performance!!
--- Interesting bugs!!


Weak/Release consistency

- New flag synchronization needed

```

A := data;           while (flag != 1) {};
membarrier;          membarrier;
flag := 1;           X := A;
    
```

- Dekker's: same as TSO
- Causal correctness provided for this code



Q: What
value will
get printed?
Answer: 1



Learning more about memory models

Shared Memory Consistency Models: A Tutorial
by Sarita Adve, Kouroush Gharachorloo
in IEEE Computer 1996

RFM: Read the F****ng Manual of the system you are working on!
(Different microprocessors and systems supports different memory models.)

Issue to think about:

What code reordering may compilers really do?
Have to use "volatile" declarations in C.

X86's current memory model

Common view in academia: TSO

If you ask Intel:

- Processor consistency with causal correctness for non-atomic memory ops
- TSO for atomic memory ops

- Video presentation:

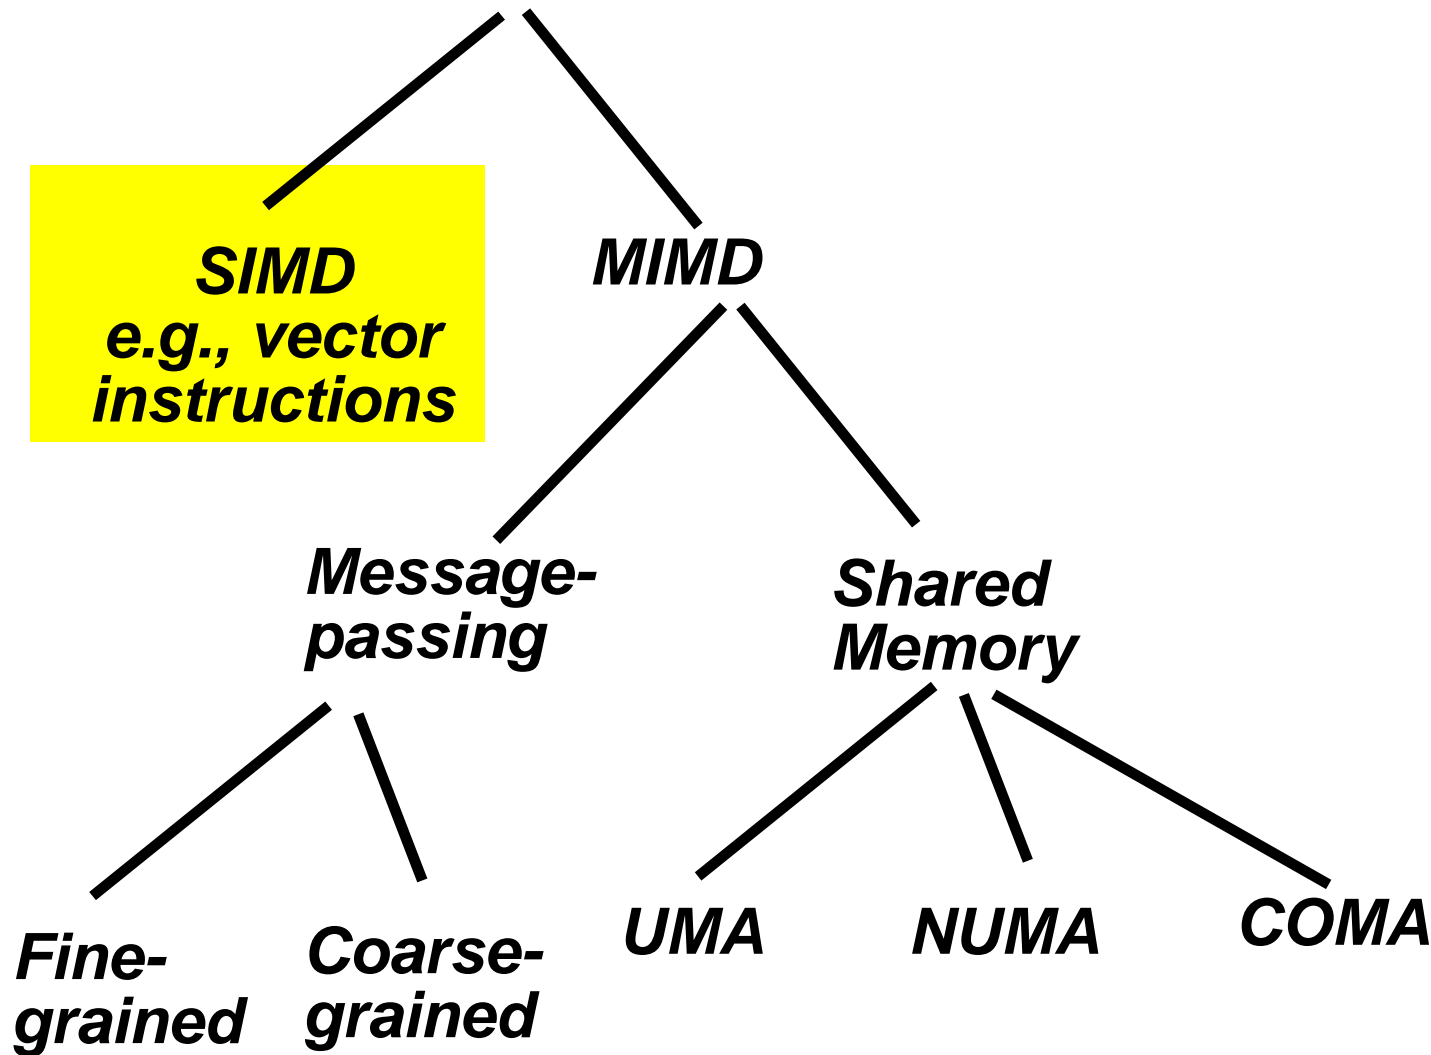
<http://www.youtube.com/watch?v=WUfvvFD5tAA&hl=sv>

- See section 8.2 in this manual:


<http://developer.intel.com/Assets/PDF/manual/253668.pdf>



A few words about SIMD



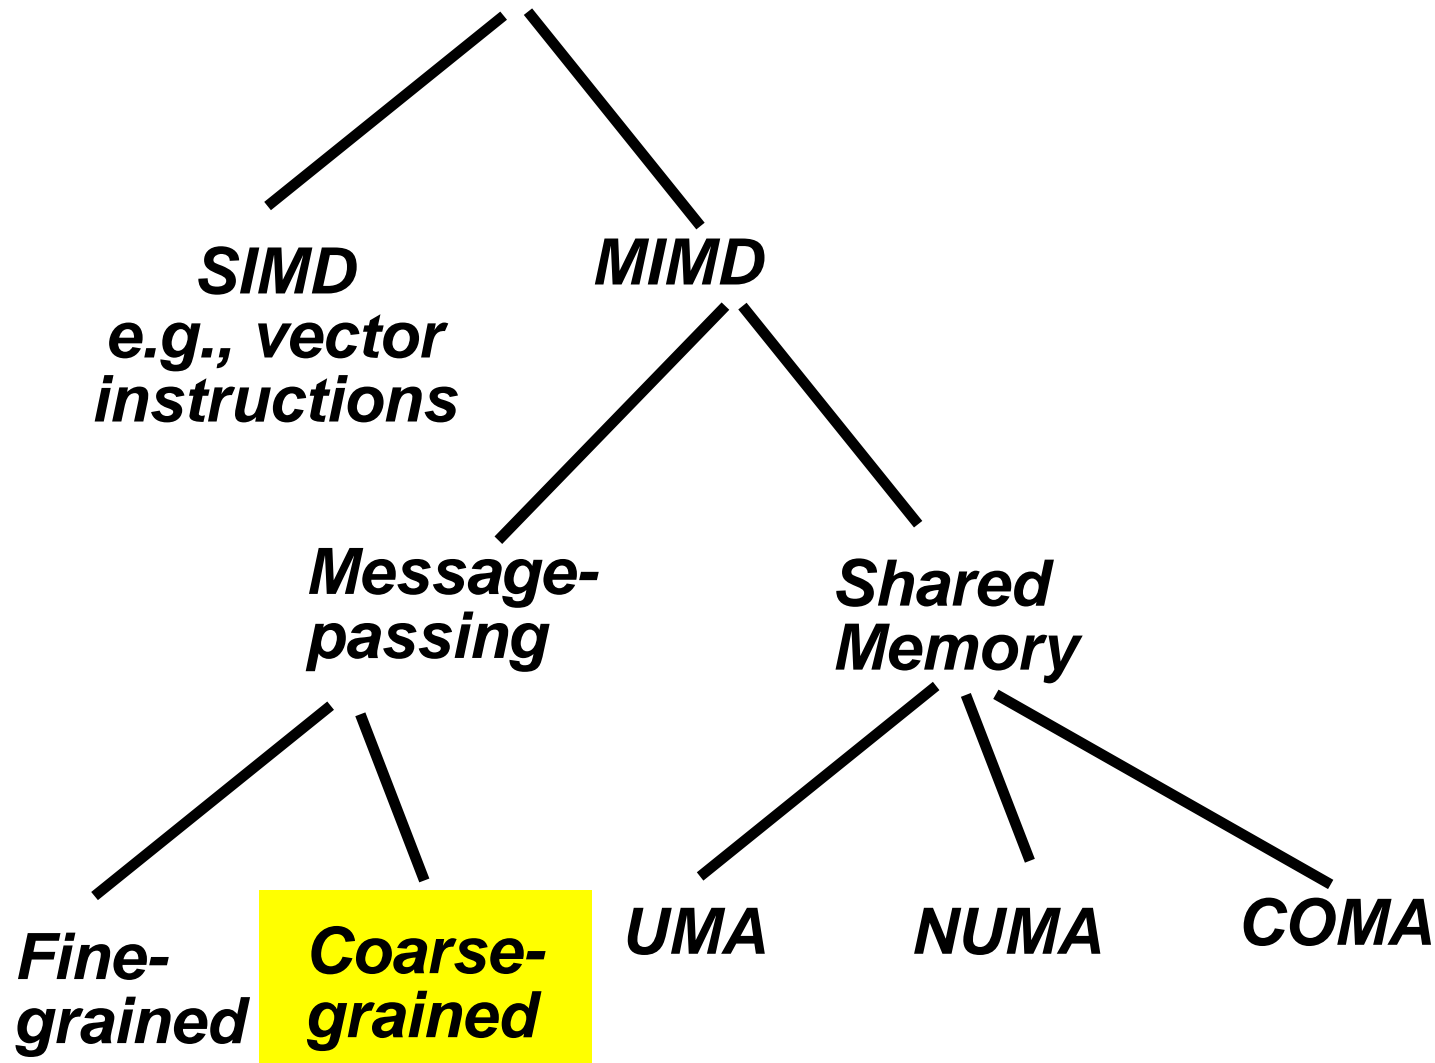
Examples of vector instructions



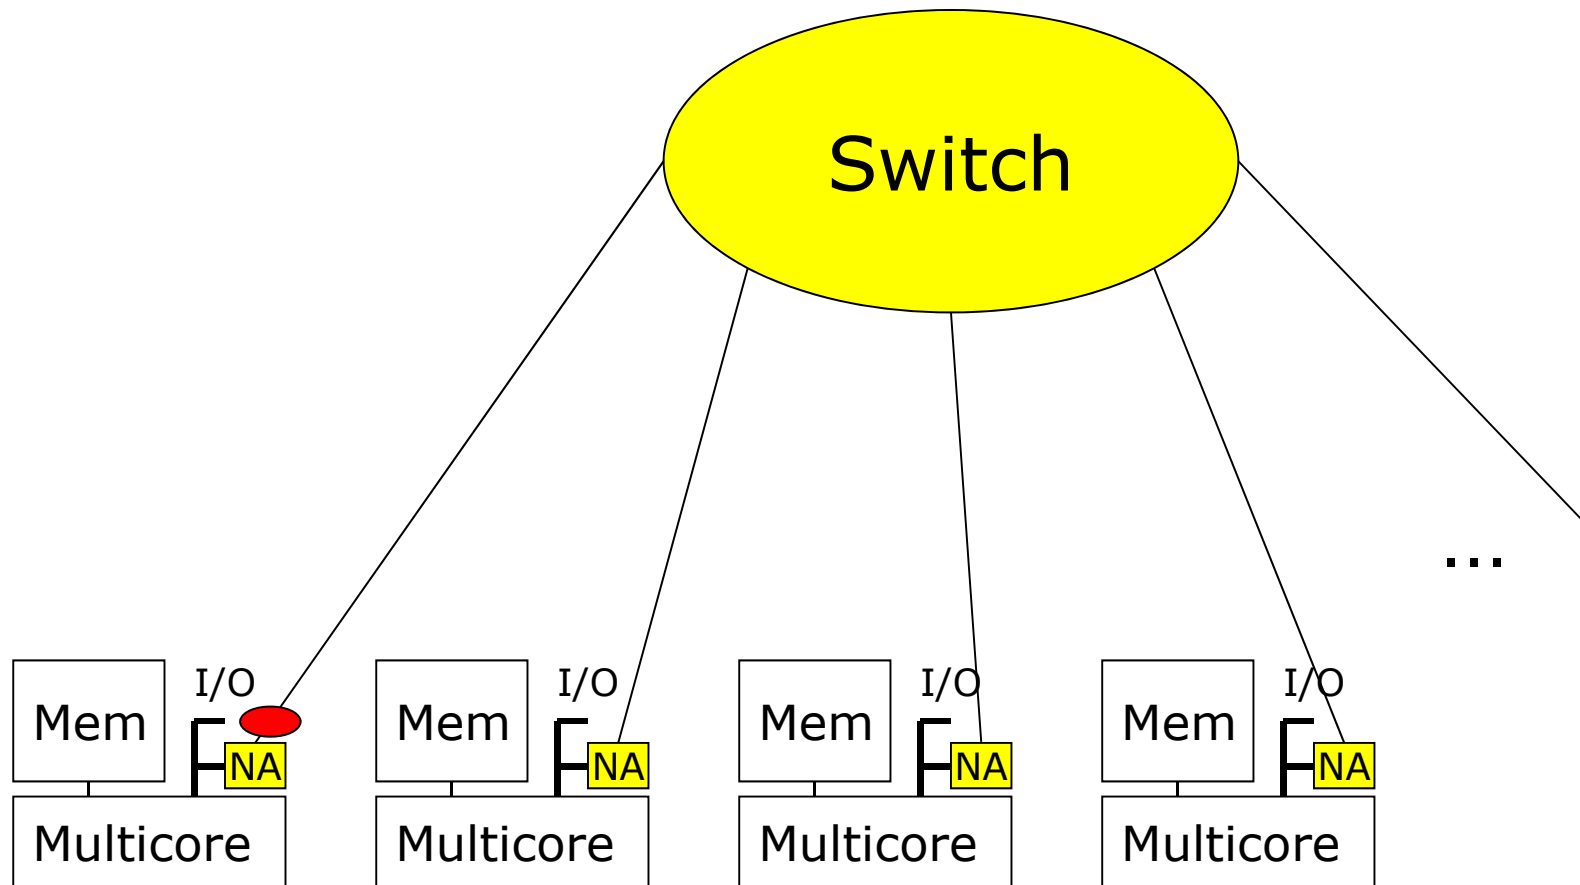
x86 Vector instructions

- MMX: 64 bit vectors (e.g., two 32bit ops)
- SSE n : 128 bit vectors(e.g., four 32 bit ops)
- AVX: 256 bit vectors(e.g., eight 32 bit ops)
(in Sandy Bridge, Q1 2011)
- Intel MIC: "16-way vectors" late 2012??

A few words on Message-passing



A modern "supercomputer"



```

X = vec[i];
MPI_send(X, to_dest);
...
  
```

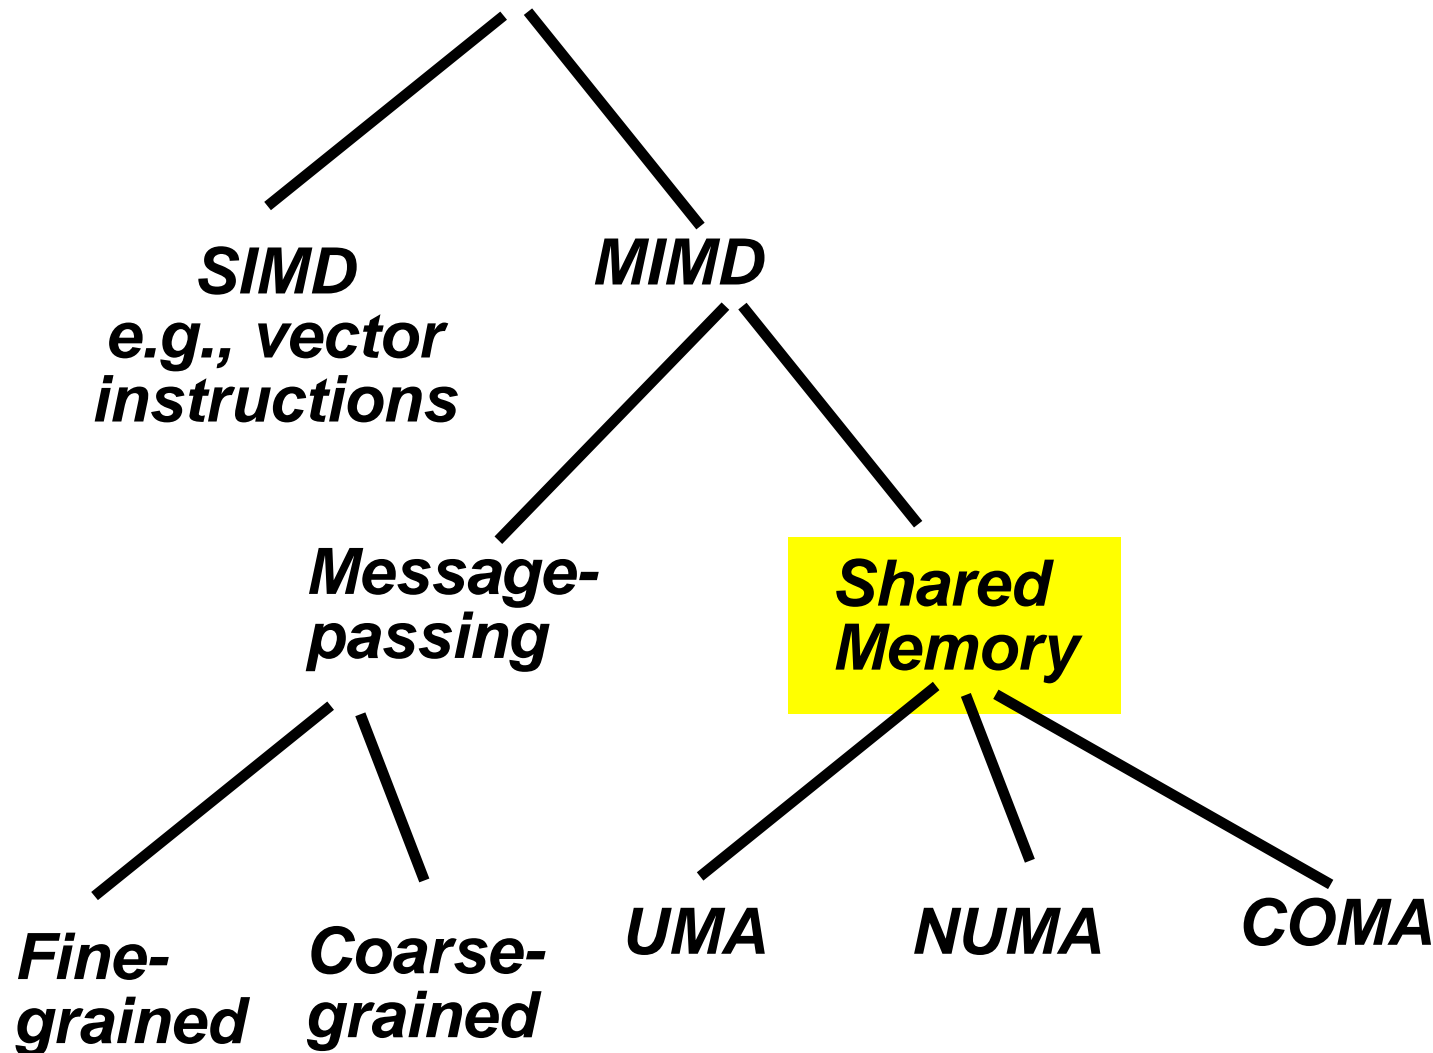
```

...
MPI_receive(Y, from_source;
print (Y);
  
```


MPI inside a multicore?

- MPI can be implemented on top of coherent shared memory
- Coherent shard memory cannot easily be implemented on top of MPI
- Many options for parallelism within a "node":
 - ✱ OpenMP
 - ✱ MPI
 - ✱ Posix threads
 - ✱ ...

A few words about simultaneously multithreading (SMT) or “Hyper-threading”



Several threads sharing a pipeline

TLP helps finding more independent instructions

